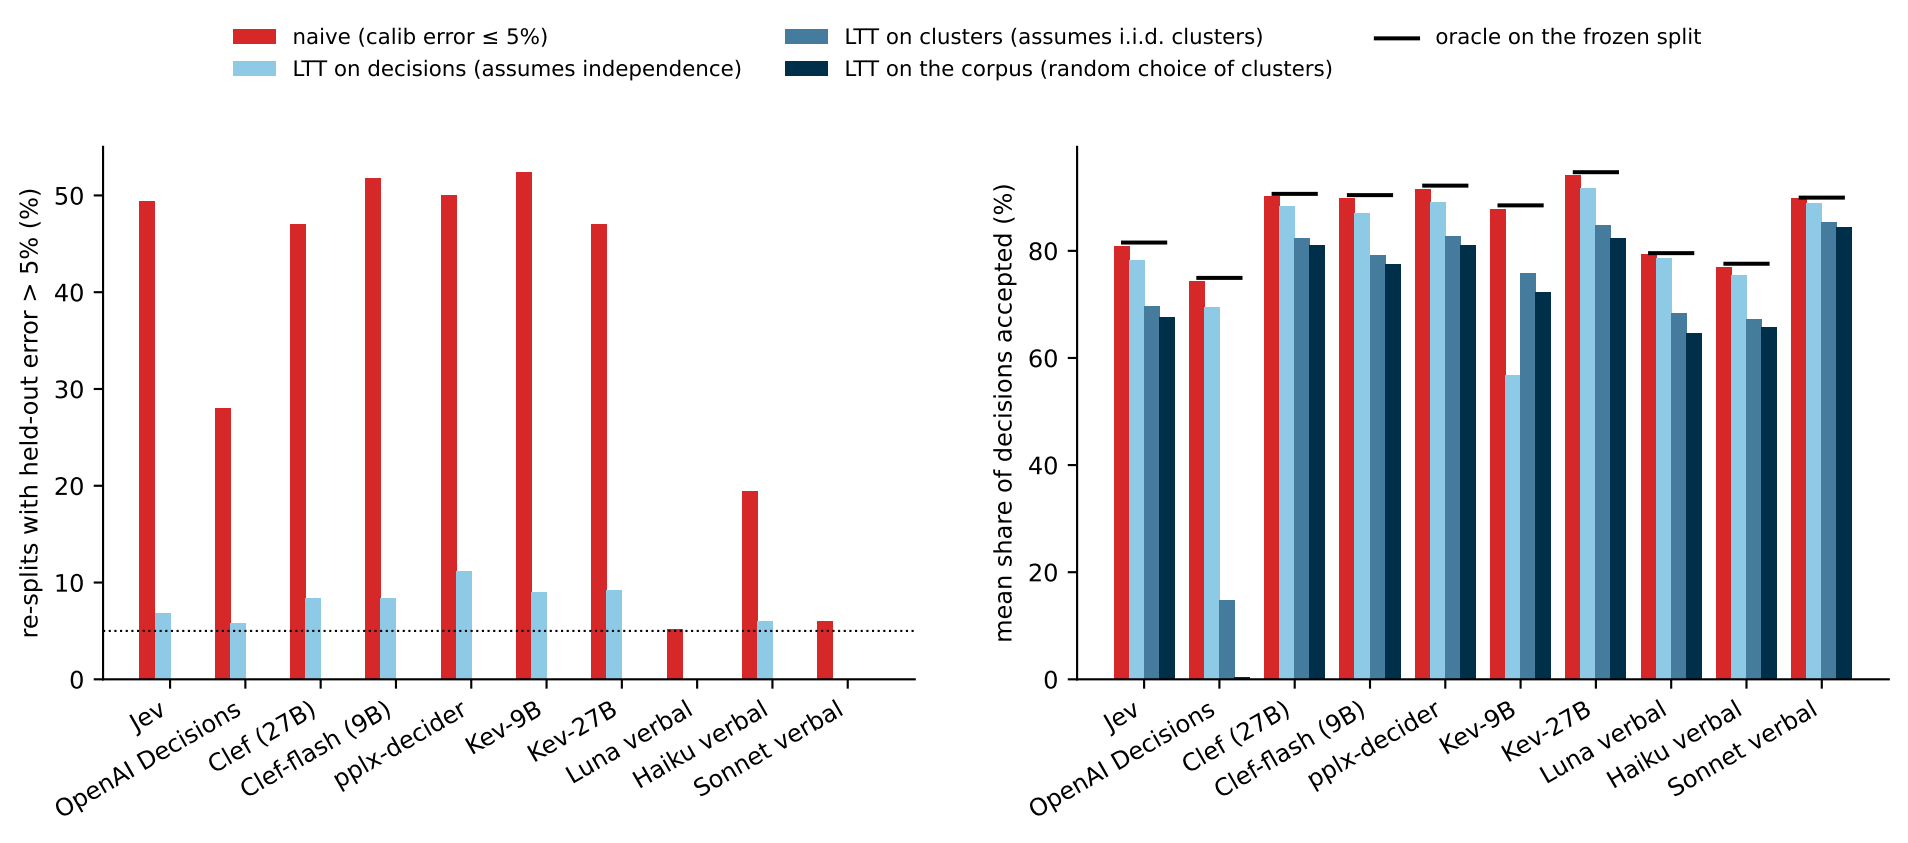
naive (calib error ≤ 5%)
LTT on decisions (assumes independence)
LTT on clusters (assumes i.i.d. clusters)
LTT on the corpus (random choice of clusters)
oracle on the frozen split
re-splits with held-out error > 5% (%)
mean share of decisions accepted (%)
0
10
20
30
40
50
0
20
40
60
80
Jev
OpenAI Decisions
Clef (27B)
Clef-flash (9B)
pplx-decider
Kev-9B
Kev-27B
Luna verbal
Haiku verbal
Sonnet verbal